

Import Required Libraries:

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader, TensorDataset
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import time
```

Device Configuration:

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)
```

Using device: cpu

Load & Preprocess Data:

```
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

train_dataset = torchvision.datasets.FashionMNIST(
    root='./data',
    train=True,
    download=True,
    transform=transform
)

100%|██████████| 26.4M/26.4M [00:01<00:00, 15.0MB/s]
100%|██████████| 29.5k/29.5k [00:00<00:00, 274kB/s]
100%|██████████| 4.42M/4.42M [00:00<00:00, 5.11MB/s]
100%|██████████| 5.15k/5.15k [00:00<00:00, 10.8MB/s]
```

Training Data:

```
train_dataset = torchvision.datasets.FashionMNIST(
    root='./data',
    train=True,
    download=True,
    transform=transform
)
```

```
train_dataset
Dataset FashionMNIST
Number of datapoints: 60000
Root location: ./data
Split: Train
StandardTransform
Transform: Compose(
    ToTensor()
    Normalize(mean=(0.5,), std=(0.5,))
)
```

Test Data:

```
df_test = pd.read_csv("/content/fashion-mnist-test.csv")

X_test = df_test.drop("label", axis=1).values.astype(np.float32)
y_test = df_test["label"].values

X_test = X_test / 255.0
X_test = (X_test - 0.5) / 0.5
```

```
X_test = X_test.reshape(-1, 1, 28, 28)

X_test = torch.tensor(X_test)
y_test = torch.tensor(y_test).long()

test_dataset = TensorDataset(X_test, y_test)
```

Data Loaders:

```
batch_size = 64

train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)
```

Multi-layer Feed Forward Neural Network:

```
class FashionNet(nn.Module):
    def __init__(self, hidden_layers, activation="relu"):
        super().__init__()

        self.layers = nn.ModuleList()
        self.activations = {}

        input_size = 28 * 28

        for i, hidden_units in enumerate(hidden_layers):
            linear = nn.Linear(input_size, hidden_units)
            self.layers.append(linear)

            if activation == "relu":
                act = nn.ReLU()
            elif activation == "sigmoid":
                act = nn.Sigmoid()
            elif activation == "tanh":
                act = nn.Tanh()
            elif activation == "leaky_relu":
                act = nn.LeakyReLU(0.01)

            self.layers.append(act)

            # Hook for visualization
            linear.register_forward_hook(self.save_activation(f"Layer_{i+1}"))
            input_size = hidden_units

        self.output = nn.Linear(input_size, 10)

    def save_activation(self, name):
        def hook(model, input, output):
            self.activations[name] = output.detach()
        return hook

    def forward(self, x):
        x = x.view(x.size(0), -1)
        for layer in self.layers:
            x = layer(x)
        return self.output(x)
```

Training Function:

```
def train_model(model, train_loader, criterion, optimizer, epochs):
    model.to(device)
    history = {"loss": [], "accuracy": []}

    for epoch in range(epochs):
        model.train()
        total_loss = 0
        correct = 0
        total = 0

        for images, labels in train_loader:
            images, labels = images.to(device), labels.to(device)
```

```

outputs = model(images)
loss = criterion(outputs, labels)

optimizer.zero_grad()
loss.backward()
optimizer.step()

total_loss += loss.item()
_, predicted = torch.max(outputs, 1)
correct += (predicted == labels).sum().item()
total += labels.size(0)

epoch_loss = total_loss / len(train_loader)
epoch_acc = 100 * correct / total

history["loss"].append(epoch_loss)
history["accuracy"].append(epoch_acc)

print(f"Epoch [{epoch+1}/{epochs}] Loss: {epoch_loss:.4f} Accuracy: {epoch_acc:.2f}%")

return history

```

Evaluation Function:

```

def evaluate_model(model, test_loader):
    model.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs, 1)
            correct += (predicted == labels).sum().item()
            total += labels.size(0)

    accuracy = 100 * correct / total
    print(f"Test Accuracy: {accuracy:.2f}%")
    return accuracy

```

Base Model Training:

```

model = FashionNet(
    hidden_layers=[128, 128, 128],
    activation="relu"
)

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

train_model(model, train_loader, criterion, optimizer, epochs=10)
evaluate_model(model, test_loader)

Epoch [1/10] Loss: 0.5183 Accuracy: 81.14%
Epoch [2/10] Loss: 0.3807 Accuracy: 86.13%
Epoch [3/10] Loss: 0.3445 Accuracy: 87.26%
Epoch [4/10] Loss: 0.3163 Accuracy: 88.28%
Epoch [5/10] Loss: 0.2971 Accuracy: 89.03%
Epoch [6/10] Loss: 0.2819 Accuracy: 89.62%
Epoch [7/10] Loss: 0.2696 Accuracy: 89.95%
Epoch [8/10] Loss: 0.2599 Accuracy: 90.29%
Epoch [9/10] Loss: 0.2495 Accuracy: 90.47%
Epoch [10/10] Loss: 0.2372 Accuracy: 91.02%
Test Accuracy: 90.32%
90.32

```

Experimenting with different depth and widths:

```

depths = [1, 3, 5]
widths = [64, 128]

for d in depths:
    for w in widths:

```

```
print(f"\nDepth: {d}, Width: {w}")
model = FashionNet([w] * d, activation="relu")
optimizer = optim.Adam(model.parameters(), lr=0.001)
train_model(model, train_loader, criterion, optimizer, epochs=5)
evaluate_model(model, test_loader)
```

Depth: 1, Width: 64
Epoch [1/5] Loss: 0.5224 Accuracy: 81.35%
Epoch [2/5] Loss: 0.4008 Accuracy: 85.56%
Epoch [3/5] Loss: 0.3632 Accuracy: 86.97%
Epoch [4/5] Loss: 0.3408 Accuracy: 87.61%
Epoch [5/5] Loss: 0.3184 Accuracy: 88.41%
Test Accuracy: 88.18%

Depth: 1, Width: 128
Epoch [1/5] Loss: 0.5033 Accuracy: 81.75%
Epoch [2/5] Loss: 0.3784 Accuracy: 86.22%
Epoch [3/5] Loss: 0.3388 Accuracy: 87.52%
Epoch [4/5] Loss: 0.3133 Accuracy: 88.42%
Epoch [5/5] Loss: 0.2978 Accuracy: 88.97%
Test Accuracy: 89.33%

Depth: 3, Width: 64
Epoch [1/5] Loss: 0.5575 Accuracy: 79.68%
Epoch [2/5] Loss: 0.3962 Accuracy: 85.45%
Epoch [3/5] Loss: 0.3554 Accuracy: 86.89%
Epoch [4/5] Loss: 0.3334 Accuracy: 87.68%
Epoch [5/5] Loss: 0.3123 Accuracy: 88.38%
Test Accuracy: 88.41%

Depth: 3, Width: 128
Epoch [1/5] Loss: 0.5188 Accuracy: 80.98%
Epoch [2/5] Loss: 0.3776 Accuracy: 86.03%
Epoch [3/5] Loss: 0.3391 Accuracy: 87.47%
Epoch [4/5] Loss: 0.3187 Accuracy: 88.17%
Epoch [5/5] Loss: 0.3002 Accuracy: 88.78%
Test Accuracy: 89.07%

Depth: 5, Width: 64
Epoch [1/5] Loss: 0.6249 Accuracy: 76.55%
Epoch [2/5] Loss: 0.4227 Accuracy: 84.46%
Epoch [3/5] Loss: 0.3818 Accuracy: 85.94%
Epoch [4/5] Loss: 0.3537 Accuracy: 86.86%
Epoch [5/5] Loss: 0.3326 Accuracy: 87.77%
Test Accuracy: 88.88%

Depth: 5, Width: 128
Epoch [1/5] Loss: 0.5856 Accuracy: 78.20%
Epoch [2/5] Loss: 0.4106 Accuracy: 84.93%
Epoch [3/5] Loss: 0.3662 Accuracy: 86.42%
Epoch [4/5] Loss: 0.3393 Accuracy: 87.49%
Epoch [5/5] Loss: 0.3197 Accuracy: 88.23%
Test Accuracy: 88.74%

Different Activation Functions:

```
activations = ["relu", "sigmoid", "tanh", "leaky_relu"]

for act in activations:
    print(f"\nActivation: {act}")
    start = time.time()

    model = FashionNet([128, 128, 128], activation=act)
    optimizer = optim.Adam(model.parameters(), lr=0.001)

    train_model(model, train_loader, criterion, optimizer, epochs=5)
    evaluate_model(model, test_loader)

    print("Training Time:", round(time.time() - start, 2), "seconds")
```

Activation: relu
Epoch [1/5] Loss: 0.5175 Accuracy: 81.09%
Epoch [2/5] Loss: 0.3740 Accuracy: 86.22%
Epoch [3/5] Loss: 0.3403 Accuracy: 87.40%
Epoch [4/5] Loss: 0.3141 Accuracy: 88.39%
Epoch [5/5] Loss: 0.2974 Accuracy: 88.97%
Test Accuracy: 88.56%
Training Time: 82.07 seconds

```
Activation: sigmoid
Epoch [1/5] Loss: 0.8696 Accuracy: 67.67%
Epoch [2/5] Loss: 0.4681 Accuracy: 83.75%
Epoch [3/5] Loss: 0.3952 Accuracy: 85.94%
Epoch [4/5] Loss: 0.3557 Accuracy: 87.40%
Epoch [5/5] Loss: 0.3332 Accuracy: 88.07%
Test Accuracy: 88.08%
Training Time: 80.44 seconds
```

```
Activation: tanh
Epoch [1/5] Loss: 0.5123 Accuracy: 81.56%
Epoch [2/5] Loss: 0.3882 Accuracy: 85.81%
Epoch [3/5] Loss: 0.3548 Accuracy: 87.05%
Epoch [4/5] Loss: 0.3364 Accuracy: 87.60%
Epoch [5/5] Loss: 0.3199 Accuracy: 88.28%
Test Accuracy: 88.30%
Training Time: 82.43 seconds
```

```
Activation: leaky_relu
Epoch [1/5] Loss: 0.5222 Accuracy: 80.65%
Epoch [2/5] Loss: 0.3822 Accuracy: 86.01%
Epoch [3/5] Loss: 0.3416 Accuracy: 87.39%
Epoch [4/5] Loss: 0.3170 Accuracy: 88.21%
Epoch [5/5] Loss: 0.2972 Accuracy: 89.01%
Test Accuracy: 89.96%
Training Time: 79.04 seconds
```

Visualization of hidden layers:

```
def visualize_activations(model, image):
    model.eval()
    with torch.no_grad():
        _ = model(image.to(device))

    plt.figure(figsize=(12, 4))
    plt.subplot(1, len(model.activations) + 1, 1)
    plt.imshow(image[0].cpu().squeeze(), cmap="gray")
    plt.title("Input")
    plt.axis("off")

    for i, (name, act) in enumerate(model.activations.items()):
        plt.subplot(1, len(model.activations) + 1, i + 2)
        plt.plot(act[0].cpu().numpy()[:100])
        plt.title(name)
        plt.axis("off")

    plt.show()
```

```
sample_image, _ = next(iter(test_loader))
visualize_activations(model, sample_image[:1])
```



